

## 03\_feature\_engineering\_text

August 21, 2025

### 0.0.1 Using LLM for Theme Extraction

In this step, we transform the raw **transaction\_description** field into structured **themes**.

- **What we will be doing now**
  - For each contract's **transaction\_description**, we send the text to an **LLM**.
  - The LLM classifies the description into one of a controlled set of themes (e.g., *IT Services*, *Construction*, *Healthcare*, *Defense Equipment*).
  - The output is stored as a new feature, **theme**, for use in downstream modeling.

#### **Example:**

- Input: "PROCUREMENT OF SERVER MAINTENANCE AND SUPPORT"
- LLM Output: "IT Services"

This converts unstructured text into a consistent categorical feature without requiring custom text preprocessing pipelines.

---

### 0.0.2 What Could Have Been Done with More Data

If we had a **very large number of lapse contracts** (positive class examples), additional strategies could have been explored to use llm (only few calls to manage cost):

- **Clustering Descriptions**
  - Apply unsupervised methods (e.g., k-means, hierarchical clustering, or embeddings + cosine similarity) on **transaction\_description**.
  - Group contracts into clusters of semantically similar descriptions before sending cluster representatives to the LLM.
  - This reduces API calls and ensures that descriptions with similar wording map to the same theme.
- **Encoding Before LLM**
  - Instead of sending raw text, first transform descriptions into vector embeddings or cluster IDs.
  - Then only send representative samples to the LLM for classification into themes.
- **Why we didn't do this here**

- Our dataset has **relatively few lapse contracts** compared to the total population.
- Over-engineering text clustering/encoding before the LLM would risk distorting the balance between **lapse** and **non-lapse** examples.
- To keep signals consistent, we directly map each contract's description to a theme using the LLM.

```
[1]: # Standard library
import os
import re
import time
import logging
from typing import List, Tuple

# Third-party Library
import pandas as pd

# LLM API
from groq import Groq
```

```
[2]: # Configure logging
logging.basicConfig(
    level=logging.INFO,
    format="%(asctime)s | %(levelname)s | %(message)s"
)
```

```
[3]: df = pd.read_csv('fea_eng_basic.csv')
df
```

```
[3]:
```

	federal_action_obligation	total_dollars_obligated \
0	23116.94	23116.94
1	84000.00	284319.29
2	66.00	66.00
3	1216.40	1216.40
4	259.70	259.70
...	...	...
2627	22973.12	119973.12
2628	1008.64	1008.64
2629	374.70	374.70
2630	289.34	289.34
2631	26.62	26.62

	current_total_value_of_award	potential_total_value_of_award \
0	23116.94	23116.94
1	284319.29	599885.84
2	66.00	66.00
3	1216.40	1216.40

4	259.70	259.70
...	...	...
2627	119973.12	119973.12
2628	1008.64	1008.64
2629	374.70	374.70
2630	289.34	289.34
2631	26.62	26.62

	action_date_fiscal_year	funding_agency_code	award_type	\
0	2022	13	1	
1	2022	36	1	
2	2022	15	1	
3	2022	97	0	
4	2022	97	0	
...	...	...	...	
2627	2023	97	1	
2628	2022	15	0	
2629	2022	15	0	
2630	2022	97	0	
2631	2022	15	0	

	type_of_contract_pricing	\
0	0	
1	0	
2	0	
3	1	
4	1	
...	...	
2627	0	
2628	0	
2629	0	
2630	1	
2631	0	

	transaction_description	extent_competed	...	\
0	WEATHER OBSERVING STATION INCLUDING INSTALLATI...	1	...	
1	RADIOPHARMACEUTICALS FOR THE MONTANA VA HEALTH...	1	...	
2	OPEN MARKET PHARMACEUTICALS ORDER	3	...	
3	4556151540!BLANKET DISP COMFORT 1	0	...	
4	4556017656!CLEANING MONITOR	0	...	
...	...	...	...	
2627	ELECTRONIC CONTROL	4	...	
2628	FY22 MCKESSON CONTROLLED CONTRACTED PHARMACEU...	0	...	
2629	PHARMACY ORDER FOR INMATES INCARCERATED AT FCI...	0	...	
2630	4557236030!FILE K 25MM #15 6S	0	...	
2631	JULY MCKESSON CONTROLLED SUBSTANCES	0	...	

	sale	revt	ib	lt	ceq	oancf \
0	39211.000	39211.000	7725.000	54146.000	40793.000	9312.000
1	40.697	40.697	-54.454	125.427	52.413	-48.746
2	583.187	583.187	7.729	400.966	667.099	-66.537
3	12401.021	12401.021	631.232	3804.587	3425.126	709.580
4	12401.021	12401.021	631.232	3804.587	3425.126	709.580
...	...	...	...	...	...	...
2627	26.074	26.074	0.481	9.036	20.020	0.250
2628	583.187	583.187	7.729	400.966	667.099	-66.537
2629	583.187	583.187	7.729	400.966	667.099	-66.537
2630	12401.021	12401.021	631.232	3804.587	3425.126	709.580
2631	583.187	583.187	7.729	400.966	667.099	-66.537

	xrd	cogs	psc_3digit_freq	lapse_flag
0	1406.000	18171.000	120352.0	1
1	85.641	115.855	635345.0	1
2	29.307	118.470	635345.0	0
3	0.000	8534.570	120352.0	0
4	0.000	8534.570	120352.0	0
...	...	...	...	...
2627	1.828	16.821	459.0	1
2628	29.307	118.470	635345.0	0
2629	29.307	118.470	635345.0	0
2630	0.000	8534.570	463837.0	0
2631	29.307	118.470	635345.0	0

[2632 rows x 34 columns]

Check few transaction\_descriptions

```
[4]: # Filter lapsed contracts
lapsed_df = df[df["lapse_flag"] == 1]

# Randomly sample 50 transaction descriptions
sample_desc = lapsed_df["transaction_description"].dropna().sample(50,
↪random_state=42)

# Print them
for i, desc in enumerate(sample_desc, 1):
    print(f"{i}. {desc}")
```

1. WIRELESS LAN INFRASTRUCTURE LCR
2. REMANUFACTURE OF FLOW CONTROL VALVE
3. MIGRATED ID08190050 DYNAMIC AND EVOLVING FEDERAL ENTERPRISE NETWORK DEFENSE GROUP E DEFEND E
4. NETSPOT RADIOPHARMACEUTICAL DELIVERY OPTION YEAR 3
5. RADIUM-223 DICHLORIDE (XOFIGO)
6. PRODUCT SUPPORT PLAN TO PROVIDE EXTENDED WARRANTY ON HARDWARE/SOFTWARE AND

SOFTWARE UPDATES. INCLUDES SITE-VISIT OF UP TO 3 DAYS FOR ON-SITE SUPPORT IF NECESSARY. SERVICE TO BE PROVIDED FOR 12 MONTHS

7. EO14042 ADDING COVID-19 CLAUSE AS REQUIRED BY EXECUTIVE ORDER
8. REAGENTS/CONSUMABLES/CONTROLS FOR COBAS 4800 TESTING SYSTEM
9. DISPOSAL, TRANSPORTATION, AND RECYCLING OF FERROUS AND NON-FERROUS SCRAP METALS
10. CHEMISTRY/ IMMUNOCHEMISTRY EQUIPMENT, CPRR REAGENTS, OPTION YEAR 1
11. MCKESSON OPEN MARKET MEDICATIONS
12. AYDIN DISPLAYS AND OPTICONN DVI EXTENDERS
13. REGULAR MEDS: FY22 (NOVEMBER 3, 2021)
14. REAGENTS
15. MIGRATED ID08190050 DYNAMIC AND EVOLVING FEDERAL ENTERPRISE NETWORK DEFENSE GROUP E DEFEND E
16. VIDAS 3 REAGENTS AND SUPPLIES
17. REGULAR MEDS: FY22 (DECEMBER 3, 2021)
18. FUNDING INCREASE PO# 546C20054
19. CIRCUIT CARD ASSEMB
20. NANOSTRING GEOMX
21. UPDATE SOW TO REV C
22. INO NITRIC OXIDE - MODIFICATION TO UPDATE POC AND WAWF ACCEPTOR.
23. AB SCIEX LLC, QTRAP 6500+ LC/MS/MS
24. POLAR ORGANIC CHEMICAL INTEGRATIVE SAMPLERS
25. VISN 20 POC WHOLE BLOOD GLUCOSE TESTING ANALYZERS BPA
26. REAGENTS
27. DYNAMIC AND EVOLVING FEDERAL ENTERPRISE NETWORK DEFENSE GROUP E DEFEND E
28. CONTRACTOR SHALL PROVIDE "OPEN MARKET" PHARMACEUTICAL SUPPLIES.
29. 39 MULTIFUNCTIONAL DEVICES (LEASE) - ADD DFARS CLAUSE 252.223-7999. EO14042.
30. AYDIN DISPLAYS AND OPTICONN DVI EXTENDERS
31. MICROBIOLOGY TESTS, LOT
32. GXP EXPLORER ENT.
33. THE PURPOSE OF THIS MODIFICATION CHANGE THE CONTRACTING OFFICER AND CLOSEOUT CONTRACT.
34. MCKESSON GENERAL MEDICATIONS 10/01/2023-11/17/2023
35. ADDED FUNDS THAT WERE ERRONEOUSLY REMOVED.
36. BLOOD CULTURE BOTTLES: MODIFICATION TO EXERCISE OPTION QUANTITY ONE
37. AUTOMATED URINALYSIS
38. ROUTINE AND NARCOTIC MEDICINES 36W79720D0001
39. ILLUMINA INC (AMBIS #1843684)
40. CPT FOR ONE (1) TOSOH G8 HBA1C HPLC ANALYZER
41. CHEMISTRY IMMUNOASSAY ANALYZER FOR BIG SPRING VA
42. ULTRIO ELITE TEST KIT
43. CELLULAR LOAD TESTER SUPPORT - EO 14042
44. FILMARRAY DIAGNOSTICS TESTING
45. COST PER TEST - IMMUNOCHEMISTRY ANALYZER
46. EMBEDDED GLOBAL POSITIONING SYSTEM/INERTIAL NAVIGATION SYSTEM (EGI) - MODERNIZED (EGI-M) ENGINEERING, MANUFACTURING AND DEVELOPMENT (EMD) PHASE
47. PON MAINTENANCE CYBERSECURE BASE YEAR
48. MICROBIOLOGY IDENTIFICATION AND SUSECPTIBILITY

49. MIGRATED ID08190050 DYNAMIC AND EVOLVING FEDERAL ENTERPRISE NETWORK DEFENSE  
GROUP E DEFEND E  
50. GN TEST KIT/AST-GN95 PN 421982

### 0.0.3 Developing the Themes

We developed the following set of themes by **randomly sampling and reviewing many transaction\_description values**, then iteratively refining categories.

This process ensured coverage across the most common spending patterns while keeping the number of categories manageable.

Having a consistent theme mapping helps:

- Reduce noise from messy free-text descriptions.
- Provide interpretable contract groupings for downstream models.
- Make LLM calls more reliable by constraining the classification set.

---

### Theme Categories

- a. Medical & Healthcare Supplies/Services
- b. IT, Software & Hardware
- c. Facilities, Utilities & Maintenance
- d. Professional & Management Services
- e. Financial & Contract Modifications
- f. Training, Education & Outreach
- g. Logistics & Support Services
- h. Closeout & Administrative Actions
- i. Defense, Aerospace & Mechanical Equipment
- j. Other

---

### 0.0.4 Use of LLM for Feature Mapping

We will leverage a Large Language Model (LLM) to classify contract descriptions into predefined themes.

#### Process Overview:

##### 1. Batching the Data

- Large datasets are split into manageable batches (e.g., 50 descriptions per batch).

## 2. Prompt Construction

- Each batch is formatted into a structured prompt with transaction descriptions and the theme list.
- The model is instructed to assign **exactly one theme per item**.

## 3. Model Inference

- The `meta-llama/llama-4-scout-17b-16e-instruct` model is used for classification.
- The output is in the strict format: `<number> <theme_letter>`.

## 4. Parsing & Updating

- The raw output is parsed into a dictionary mapping row numbers → theme.
- The dataframe is updated with theme classifications for each transaction.

## 5. Iterative Processing

- The process repeats across all batches, with delays between requests to respect API rate limits.

## 6. Final Output

- The dataframe contains a new **Theme** column, ensuring each contract is mapped consistently.

---

### LLM Model Discussion:

- `openai/gpt-oss-120b` → Did not provide correct classification format.
- `openai/gpt-oss-20b` → Also failed to return responses in the expected structured style.
- `qwen/qwen3-32b` → Misunderstood the prompt and gave verbose answers such as:  
*“Theme a is Medical & Healthcare Supplies/Services. That would include things like medications, medical devices, healthcare services, etc. So items with drugs, medical equipment, or services related to healthcare should go here.”*
- `llama-4-scout-17b-16e-instruct` was the **only model** that consistently followed the prompt and returned clean, structured responses as required.

```
[5]: # Theme map
THEME_MAP = {
    "a": "Medical & Healthcare Supplies/Services",
    "b": "IT, Software & Hardware",
    "c": "Facilities, Utilities & Maintenance",
    "d": "Professional & Management Services",
    "e": "Financial & Contract Modifications",
    "f": "Training, Education & Outreach",
    "g": "Logistics & Support Services",
    "h": "Closeout & Administrative Actions",
    "i": "Defense, Aerospace & Mechanical Equipment",
    "j": "Other",
}

# tolerant parser: "1) a", "1) a - extra", "1) A" all OK
ROW_RE = re.compile(r"^\s*(\d+)\)\s*([a-j])\b", re.IGNORECASE)
```

```
[11]: class GroqThemeClassifier:
    """
    Classify contract descriptions into predefined themes using Groq LLM API.
    """

    def __init__(
        self,
        api_key: str = None,
        model: str = "meta-llama/llama-4-scout-17b-16e-instruct",
        rpm_delay_sec: float = 2.0,
    ):
        """
        Parameters
        -----
        api_key : str, optional
            Groq API key. If None, will read from env variable GROQ_API_KEY.
        model : str
            Groq model to use for classification.
        rpm_delay_sec : float
            Delay (seconds) between batches to respect rate limits.
        """
        self.api_key = api_key or os.getenv("GROQ_API_KEY")
        self.client = Groq(api_key=self.api_key)
        self.model = model
        self.rpm_delay_sec = rpm_delay_sec

        self.logger = logging.getLogger(self.__class__.__name__)

    @staticmethod
    def parse_mapping(text: str) -> dict:
        """Parse model output into {row_number: theme} mapping."""
        mapping = {}
        for line in str(text).splitlines():
            m = ROW_RE.match(line.strip())
            if m:
                num = int(m.group(1))
                letter = m.group(2).lower()
                mapping[num] = THEME_MAP[letter]
        return mapping

    @staticmethod
    def build_prompt(numbered_items: List[Tuple[int, str]]) -> str:
        """Build classification prompt for Groq."""
        header = [
            "Classify EACH transaction into EXACTLY ONE theme letter (a-j).",
            "",
            "Themes:",

```



```

        *[f"{k}. {v}" for k, v in THEME_MAP.items()],
        "",
        "Return ONLY lines in the exact format: <number>) <letter>",
        "Do NOT include any other text before or after.",
        "",
        "Items:",
    ]
    lines = [f"{i}) {txt}" for i, txt in numbered_items]
    return "\n".join(header + lines)

def classify_texts(self, texts: List[str], max_tokens: int = 200) -> str:
    """Send a batch of texts to Groq and return raw model output."""
    prompt = self.build_prompt([(i + 1, t) for i, t in enumerate(texts)])
    resp = self.client.chat.completions.create(
        model=self.model,
        messages=[
            {"role": "system", "content": "You are a precise classifier."},
            {"role": "user", "content": prompt},
        ],
        temperature=0,
        max_tokens=max_tokens,
    )
    return resp.choices[0].message.content

def classify_in_batches(
    self,
    df: pd.DataFrame,
    text_col: str = "transaction_description",
    theme_col: str = "Theme",
    batch_size: int = 50,
    max_batches: int = None,
    show_prompts: bool = False,
    show_model_output: bool = True,
) -> pd.DataFrame:
    """
    Classify descriptions into themes in batches.

    Parameters
    -----
    df : pd.DataFrame
        Input dataframe with a text column.
    text_col : str
        Column containing contract descriptions.
    theme_col : str
        Column to write theme classifications into.
    batch_size : int
        Number of rows per batch.
    """

```

```

max_batches : int, optional
    If set, only process up to this many batches.
show_prompts : bool
    Whether to log full prompts for debugging.
show_model_output : bool
    Whether to log raw model output.

Returns
-----
pd.DataFrame
    Dataframe with new theme assignments.
"""
if theme_col not in df.columns:
    df[theme_col] = pd.Series(index=df.index, dtype="object")

idxs = df.index[df[text_col].notna()].tolist()
total = len(idxs)
if total == 0:
    self.logger.warning("No rows to classify.")
    return df

num_batches = (total + batch_size - 1) // batch_size
if max_batches is not None:
    num_batches = min(num_batches, max_batches)

self.logger.info("Planned: %s batch(es) of up to %s rows", num_batches,
↳batch_size)

for b in range(num_batches):
    start, end = b * batch_size, min((b + 1) * batch_size, total)
    batch_index = idxs[start:end]

    numbered_items = [(i, str(df.at[idx, text_col]).strip())
        for i, idx in enumerate(batch_index, start=1)]

    prompt = self.build_prompt(numbered_items)
    if show_prompts:
        self.logger.debug("Prompt (batch %s/%s):\n%s", b + 1,
↳num_batches, prompt)

    # Call Groq API
    raw_output = self.classify_texts(
        [t for _, t in numbered_items],
        max_tokens=max(600, 10 * len(numbered_items)),
    )
    if show_model_output:

```

```

        self.logger.debug("Raw model output (batch %s/%s):\n%s", b + 1,
↪num_batches, raw_output)

        # Parse and update
        mapping = self.parse_mapping(raw_output)
        if not mapping:
            self.logger.warning("No lines parsed for batch %s. Skipping
↪updates.", b + 1)
        else:
            for local_num, idx in enumerate(batch_index, start=1):
                df.at[idx, theme_col] = mapping.get(local_num,
↪THEME_MAP["j"])

        # Rate limiting
        if b < num_batches - 1 and self.rpm_delay_sec:
            time.sleep(self.rpm_delay_sec)

        self.logger.info("Classification complete. Updated column '%s'.",
↪theme_col)
        return df

```

```

[16]: # Logging at `DEBUG` level is enabled so raw model outputs can be monitored in
↪test runs
logging.getLogger("GroqThemeClassifier").setLevel(logging.DEBUG)

```

```

[12]: # Test API Call
clf = GroqThemeClassifier()

# Test API Call
df_out = clf.classify_in_batches(
    df,
    text_col="transaction_description",
    theme_col="Theme",
    batch_size=3,          # 3 rows per batch
    max_batches=2,        # only call 2 batches
    show_prompts=True,    # log full prompts for debugging
    show_model_output=True # log raw LLM responses
)

```

2025-08-21 18:17:15,455 | INFO | Planned: 2 batch(es) of up to 3 rows  
2025-08-21 18:17:15,456 | DEBUG | Prompt (batch 1/2):  
Classify EACH transaction into EXACTLY ONE theme letter (a-j).

Themes:

- a. Medical & Healthcare Supplies/Services
- b. IT, Software & Hardware
- c. Facilities, Utilities & Maintenance

- d. Professional & Management Services
- e. Financial & Contract Modifications
- f. Training, Education & Outreach
- g. Logistics & Support Services
- h. Closeout & Administrative Actions
- i. Defense, Aerospace & Mechanical Equipment
- j. Other

Return ONLY lines in the exact format: <number>) <letter>  
Do NOT include any other text before or after.

Items:

- 1) WEATHER OBSERVING STATION INCLUDING INSTALLATION, TRAINING AND WARRANTY
  - 2) RADIOPHARMACEUTICALS FOR THE MONTANA VA HEALTH CARE SYSTEM
  - 3) OPEN MARKET PHARMACEUTICALS ORDER
- 2025-08-21 18:17:16,176 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:17:16,178 | DEBUG | Raw model output (batch 1/2):
- 1) i
  - 2) a
  - 3) a
- 2025-08-21 18:17:18,184 | DEBUG | Prompt (batch 2/2):  
Classify EACH transaction into EXACTLY ONE theme letter (a-j).

Themes:

- a. Medical & Healthcare Supplies/Services
- b. IT, Software & Hardware
- c. Facilities, Utilities & Maintenance
- d. Professional & Management Services
- e. Financial & Contract Modifications
- f. Training, Education & Outreach
- g. Logistics & Support Services
- h. Closeout & Administrative Actions
- i. Defense, Aerospace & Mechanical Equipment
- j. Other

Return ONLY lines in the exact format: <number>) <letter>  
Do NOT include any other text before or after.

Items:

- 1) 4556151540!BLANKET DISP COMFORT 1
  - 2) 4556017656!CLEANING MONITOR
  - 3) MONTHLY LEASE CATEGORY I, II, III MFDS - ADD DFARS CLAUSE 252.223-7999. EO 14042.
- 2025-08-21 18:17:18,795 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:17:18,798 | DEBUG | Raw model output (batch 2/2):
- 1) j

2) c

3) e

2025-08-21 18:17:18,799 | INFO | Classification complete. Updated column 'Theme'.

**Note:** - The GroqThemeClassifier is now working successfully on test batches. - We will now run the classification process on the **entire dataset**.

```
[15]: # Run on entire dataset
classifier = GroqThemeClassifier()

# Final dataset
df_final = clf.classify_in_batches(
    df,
    text_col="transaction_description",
    theme_col="Theme",
    batch_size=50,
    max_batches=None,
    show_prompts=False,
    show_model_output=False
)
```

2025-08-21 18:22:07,802 | INFO | Planned: 53 batch(es) of up to 50 rows

2025-08-21 18:22:09,916 | INFO | HTTP Request: POST

https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

2025-08-21 18:22:12,756 | INFO | HTTP Request: POST

https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

2025-08-21 18:22:15,556 | INFO | HTTP Request: POST

https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

2025-08-21 18:22:18,428 | INFO | HTTP Request: POST

https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

2025-08-21 18:22:21,221 | INFO | HTTP Request: POST

https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

2025-08-21 18:22:24,118 | INFO | HTTP Request: POST

https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

2025-08-21 18:22:27,136 | INFO | HTTP Request: POST

https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

2025-08-21 18:22:29,996 | INFO | HTTP Request: POST

https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

2025-08-21 18:22:32,866 | INFO | HTTP Request: POST

https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

2025-08-21 18:22:35,790 | INFO | HTTP Request: POST

https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

2025-08-21 18:22:38,714 | INFO | HTTP Request: POST

https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

2025-08-21 18:22:41,777 | INFO | HTTP Request: POST

https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

2025-08-21 18:22:44,617 | INFO | HTTP Request: POST

https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

2025-08-21 18:22:47,509 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:22:50,375 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:22:53,218 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:22:55,990 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:22:59,334 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:02,368 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:05,337 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:08,138 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:10,958 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:13,823 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:16,605 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:19,482 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:22,356 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:25,152 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:27,956 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:30,826 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:33,663 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:36,599 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:39,395 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:42,166 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:44,951 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:47,820 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:50,689 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"  
2025-08-21 18:23:53,546 | INFO | HTTP Request: POST  
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"

```

2025-08-21 18:23:56,379 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:23:59,168 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:01,958 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:04,818 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:07,997 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:10,862 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:13,660 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:16,476 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:19,357 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:22,490 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:25,281 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:28,170 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:31,030 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:33,889 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:36,853 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:39,474 | INFO | HTTP Request: POST
https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 200 OK"
2025-08-21 18:24:39,478 | INFO | Classification complete. Updated column
'Theme'.

```

```
[17]: df_final
```

```

[17]:      federal_action_obligation  total_dollars_obligated \
0                23116.94                23116.94
1                84000.00                284319.29
2                 66.00                 66.00
3                1216.40                1216.40
4                 259.70                 259.70
...                ...                ...
2627            22973.12            119973.12
2628             1008.64             1008.64
2629             374.70             374.70

```

2630	289.34	289.34
2631	26.62	26.62

	current_total_value_of_award	potential_total_value_of_award	\
0	23116.94	23116.94	
1	284319.29	599885.84	
2	66.00	66.00	
3	1216.40	1216.40	
4	259.70	259.70	
...	...	...	
2627	119973.12	119973.12	
2628	1008.64	1008.64	
2629	374.70	374.70	
2630	289.34	289.34	
2631	26.62	26.62	

	action_date_fiscal_year	funding_agency_code	award_type	\
0	2022	13	1	
1	2022	36	1	
2	2022	15	1	
3	2022	97	0	
4	2022	97	0	
...	...	...	...	
2627	2023	97	1	
2628	2022	15	0	
2629	2022	15	0	
2630	2022	97	0	
2631	2022	15	0	

	type_of_contract_pricing	\
0	0	
1	0	
2	0	
3	1	
4	1	
...	...	
2627	0	
2628	0	
2629	0	
2630	1	
2631	0	

	transaction_description	extent_competed	...	\
0	WEATHER OBSERVING STATION INCLUDING INSTALLATI...	1	...	
1	RADIOPHARMACEUTICALS FOR THE MONTANA VA HEALTH...	1	...	
2	OPEN MARKET PHARMACEUTICALS ORDER	3	...	
3	4556151540!BLANKET DISP COMFORT 1	0	...	



4	4556017656!CLEANING MONITOR	0	...
...	...	...	...
2627	ELECTRONIC CONTROL	4	...
2628	FY22 MCKESSON CONTROLLED CONTRACTED PHARMACEU...	0	...
2629	PHARMACY ORDER FOR INMATES INCARCERATED AT FCI...	0	...
2630	4557236030!FILE K 25MM #15 6S	0	...
2631	JULY MCKESSON CONTROLLED SUBSTANCES	0	...

	revt	ib	lt	ceq	oancf	xrd	\
0	39211.000	7725.000	54146.000	40793.000	9312.000	1406.000	
1	40.697	-54.454	125.427	52.413	-48.746	85.641	
2	583.187	7.729	400.966	667.099	-66.537	29.307	
3	12401.021	631.232	3804.587	3425.126	709.580	0.000	
4	12401.021	631.232	3804.587	3425.126	709.580	0.000	
...	...	...	...	...	...	...	
2627	26.074	0.481	9.036	20.020	0.250	1.828	
2628	583.187	7.729	400.966	667.099	-66.537	29.307	
2629	583.187	7.729	400.966	667.099	-66.537	29.307	
2630	12401.021	631.232	3804.587	3425.126	709.580	0.000	
2631	583.187	7.729	400.966	667.099	-66.537	29.307	

	cogs	psc_3digit_freq	lapse_flag	\
0	18171.000	120352.0	1	
1	115.855	635345.0	1	
2	118.470	635345.0	0	
3	8534.570	120352.0	0	
4	8534.570	120352.0	0	
...	...	...	...	
2627	16.821	459.0	1	
2628	118.470	635345.0	0	
2629	118.470	635345.0	0	
2630	8534.570	463837.0	0	
2631	118.470	635345.0	0	

	Theme
0	Defense, Aerospace & Mechanical Equipment
1	Medical & Healthcare Supplies/Services
2	Medical & Healthcare Supplies/Services
3	Facilities, Utilities & Maintenance
4	Facilities, Utilities & Maintenance
...	...
2627	IT, Software & Hardware
2628	Financial & Contract Modifications
2629	Medical & Healthcare Supplies/Services
2630	Defense, Aerospace & Mechanical Equipment
2631	Medical & Healthcare Supplies/Services

[2632 rows x 35 columns]

```
[18]: # Check if any Theme is NaN
df_final['Theme'].isnull().sum()
```

```
[18]: np.int64(0)
```

```
[20]: # Map each unique theme label to a numeric code.
unique_sorted = sorted(set(df_final['Theme'].dropna()))
theme_to_code_alpha = {t: i for i, t in enumerate(unique_sorted)}
df_final['ThemeCodeAlpha'] = df_final['Theme'].map(theme_to_code_alpha)
```

```
[23]: # No need of string column now
df_final = df_final.drop(columns=["transaction_description", "Theme"],
    ↪errors="ignore")
```

```
[24]: # Check if all columns are numeric
all_numeric = df_final.apply(pd.api.types.is_numeric_dtype).all()
print("All columns numeric:", all_numeric)
```

All columns numeric: True

```
[25]: list(df_final.columns)
```

```
[25]: ['federal_action_obligation',
      'total_dollars_obligated',
      'current_total_value_of_award',
      'potential_total_value_of_award',
      'action_date_fiscal_year',
      'funding_agency_code',
      'award_type',
      'type_of_contract_pricing',
      'extent_competed',
      'government_furnished_property',
      'undefinitized_action',
      'performance_based_service_acquisition',
      'veteran_owned_business',
      'woman_owned_business',
      'minority_owned_business',
      'contracting_officers_determination_of_business_size',
      'foreign_owned',
      'for_profit_organization',
      'nonprofit_organization',
      'the_ability_one_program',
      'small_disadvantaged_business',
      'sic4',
      'at',
      'sale',
```

```

'revt',
'ib',
'lt',
'ceq',
'oancf',
'xrd',
'cogs',
'psc_3digit_freq',
'lapse_flag',
'ThemeCodeAlpha']

```

```

[26]: df_final.to_csv('final_feature_eng.csv', index=False)
      #df_final = pd.read_csv('final_feature_eng.csv')

```

### 0.0.5 Comparing Performance

We first trained simple models **before including ThemeCodeAlpha** as a feature. And we had:

#### Random Forest

- Accuracy: **0.8221**
- ROC-AUC: **0.824**
- Precision-Recall AUC: **0.771**

#### XGBoost

- Accuracy: **0.8322**
- ROC-AUC: **0.883**
- Precision-Recall AUC: **0.883**

```

[43]: # Import Libraries
      from sklearn.svm import SVC
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score
      from xgboost import XGBClassifier
      from lightgbm import LGBMClassifier
      from catboost import CatBoostClassifier
      from sklearn.metrics import roc_curve, roc_auc_score, precision_recall_curve,
      ↪auc

      import matplotlib.pyplot as plt

      # Ignore warning
      import warnings
      warnings.filterwarnings("ignore")

```

```

[44]: # Train = 2022 + 2023
      train_df = df_final[df_final["action_date_fiscal_year"].isin([2022, 2023])].
      ↪copy()

```

```

# Test = 2024
test_df = df_final[df_final["action_date_fiscal_year"] == 2024].copy()

# Remove date column
train_df = train_df.drop(columns="action_date_fiscal_year", errors="ignore")
test_df = test_df.drop(columns="action_date_fiscal_year", errors="ignore")

```

```

[45]: # --- Separate features and target ---
X_train = train_df.drop(columns=["lapse_flag"])
y_train = train_df["lapse_flag"]

X_test = test_df.drop(columns=["lapse_flag"])
y_test = test_df["lapse_flag"]

```

```

[46]: # 1. Linear Regression
lr = LinearRegression()
lr.fit(X_train, y_train)
y_pred_lr = (lr.predict(X_test) > 0.5).astype(int)
print("Linear Regression Accuracy:", accuracy_score(y_test, y_pred_lr))

# 2. Logistic Regression
logr = LogisticRegression(max_iter=1000, random_state=42)
logr.fit(X_train, y_train)
y_pred_logr = logr.predict(X_test)
print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred_logr))

# 3. Support Vector Machine (SVM)
svm = SVC(kernel="rbf", probability=True, random_state=42)
svm.fit(X_train, y_train)
y_pred_svm = svm.predict(X_test)
print("SVM Accuracy:", accuracy_score(y_test, y_pred_svm))

# 4. Random Forest
rf = RandomForestClassifier(n_estimators=200, random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))

# 5. XGBoost
xgb = XGBClassifier(
    use_label_encoder=False, eval_metric="logloss", random_state=42
)
xgb.fit(X_train, y_train)
y_pred_xgb = xgb.predict(X_test)
print("XGBoost Accuracy:", accuracy_score(y_test, y_pred_xgb))

# 6. LightGBM

```

```

lgbm = LGBMClassifier(n_estimators=200, random_state=42)
lgbm.fit(X_train, y_train)
y_pred_lgbm = lgbm.predict(X_test)
print("LightGBM Accuracy:", accuracy_score(y_test, y_pred_lgbm))

# 7. CatBoost
# (silent=True suppresses training logs)
cat = CatBoostClassifier(iterations=200, random_state=42, verbose=False)
cat.fit(X_train, y_train)
y_pred_cat = cat.predict(X_test)
print("CatBoost Accuracy:", accuracy_score(y_test, y_pred_cat))

```

Linear Regression Accuracy: 0.7013422818791947  
 Logistic Regression Accuracy: 0.6845637583892618  
 SVM Accuracy: 0.5100671140939598  
 Random Forest Accuracy: 0.825503355704698  
 XGBoost Accuracy: 0.8322147651006712  
 [LightGBM] [Info] Number of positive: 1167, number of negative: 1167  
 [LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000488 seconds.  
 You can set `force\_row\_wise=true` to remove the overhead.  
 And if memory is not enough, you can set `force\_col\_wise=true`.  
 [LightGBM] [Info] Total Bins 2389  
 [LightGBM] [Info] Number of data points in the train set: 2334, number of used features: 28  
 [LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 -> initscore=0.000000  
 LightGBM Accuracy: 0.8355704697986577  
 CatBoost Accuracy: 0.8322147651006712

```

[47]: models = {
    "Random Forest": rf,
    "XGBoost": xgb,
    "LightGBM": lgbm,
    "CatBoost": cat,
}

# --- ROC Curves ---
plt.figure(figsize=(8,6))
for name, model in models.items():
    if hasattr(model, "predict_proba"):
        y_proba = model.predict_proba(X_test)[: , 1]
    elif hasattr(model, "decision_function"):
        y_proba = model.decision_function(X_test)
    else:
        y_proba = model.predict(X_test)

# ROC

```

```

fpr, tpr, _ = roc_curve(y_test, y_proba)
auc_score = roc_auc_score(y_test, y_proba)
print(f"{name} ROC AUC: {auc_score:.3f}")
plt.plot(fpr, tpr, label=f"{name} (AUC={auc_score:.3f})")

plt.plot([0,1], [0,1], "k--")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curves")
plt.legend()
plt.show()

# --- Precision-Recall Curves ---
plt.figure(figsize=(8,6))
for name, model in models.items():
    if hasattr(model, "predict_proba"):
        y_proba = model.predict_proba(X_test)[: , 1]
    elif hasattr(model, "decision_function"):
        y_proba = model.decision_function(X_test)
    else:
        y_proba = model.predict(X_test)

    precision, recall, _ = precision_recall_curve(y_test, y_proba)
    pr_auc = auc(recall, precision)
    print(f"{name} Precision-Recall AUC: {pr_auc:.3f}")
    plt.plot(recall, precision, label=f"{name} (AUC={pr_auc:.3f})")

plt.xlabel("Recall")
plt.ylabel("Precision")
plt.title("Precision-Recall Curves")
plt.legend()
plt.show()


```

Random Forest ROC AUC: 0.817

XGBoost ROC AUC: 0.898

LightGBM ROC AUC: 0.889

CatBoost ROC AUC: 0.854

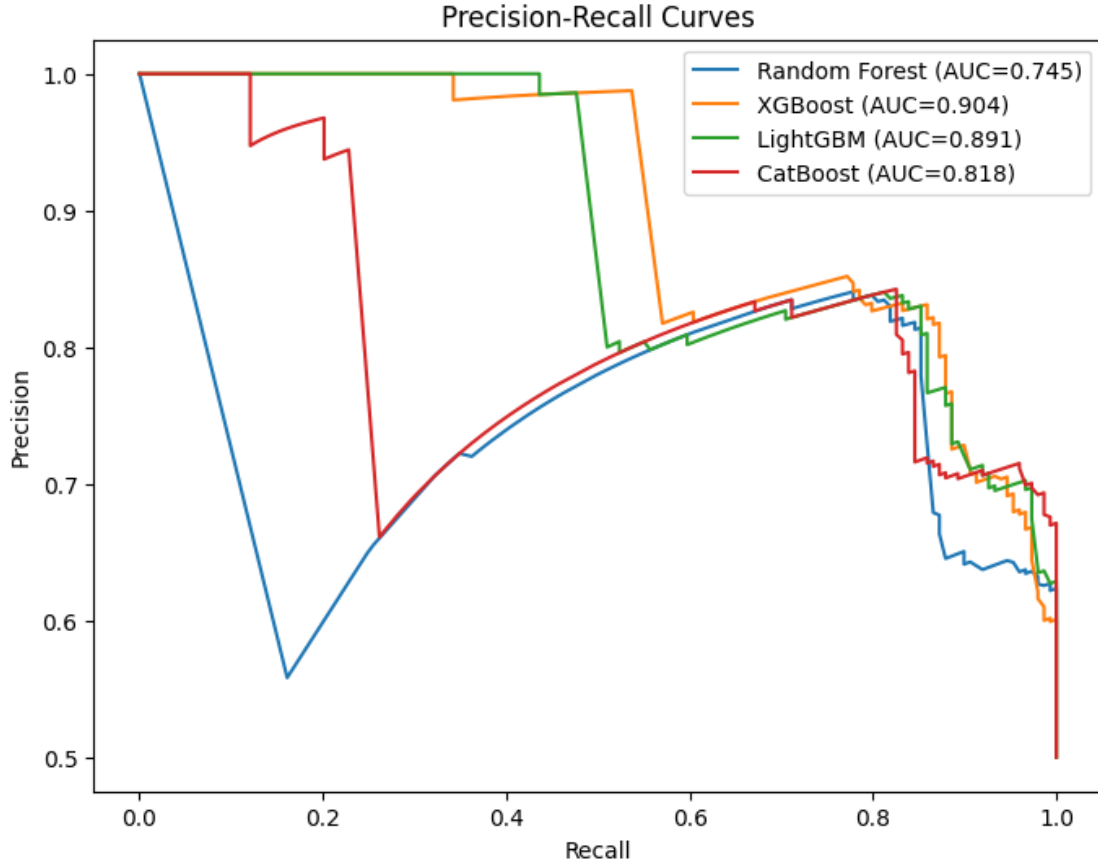


Random Forest Precision-Recall AUC: 0.745

XGBoost Precision-Recall AUC: 0.904

LightGBM Precision-Recall AUC: 0.891

CatBoost Precision-Recall AUC: 0.818



### 0.1 Comparing Model Performance

After removing the introducing the **ThemeCodeAlpha** mapping from LLM classification, we compared performance with our earlier results.

### 0.2 Model Performance Comparison

Model	Metric	Previous (w/o ThemeCodeAlpha)	Current (with ThemeCodeAlpha)
<b>Random Forest</b>	Accuracy	0.8221	0.8255
	ROC-AUC	0.824	0.817
	Precision-Recall AUC	0.771	0.745
<b>XGBoost</b>	Accuracy	0.8322	0.8322
	ROC-AUC	0.883	0.898
	Precision-Recall AUC	0.883	0.904



---

### 0.2.1 Discussion

- Random Forest showed a **slight gain in accuracy**, but **decrease in ROC-AUC and PR-AUC**, indicating weaker probability calibration and ranking ability.
- XGBoost maintained **competitive accuracy** and improved on both **ROC-AUC** (0.898 vs 0.883) and **Precision-Recall AUC** (0.904 vs 0.883).
- This suggests that **ThemeCodeAlpha** provides little meaningful additional signal.

---

**Next Step** We will proceed with **fine-tuning XGBoost**, as it achieves the **highest AUC scores** while maintaining **competitive accuracy**, making it the best candidate for modeling contract lapse risk.